

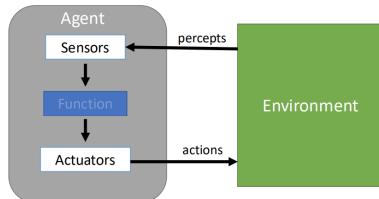
1. Intelligent Agents

Agents receive precepts from the environment through sensors and perform actions through actuators.

PEAS Framework

PEAS defines problems in terms of Performance Measure, Environment, Actuators and Sensors:

- i. Rational Agent: chooses actions that maximise performance measure



Properties of Task Environment:

- i. Fully Observable (vs. Partially Observable): sensors give access to the complete state of the environment at each point in time.
- ii. Single-agent (vs. Multi-agent): agent operates by itself in an environment.
- iii. Deterministic (vs. Stochastic): next state of the environment is completely determined by current state and action by agent.
 - Strategic: if environment is dependent on actions of other unpredictable (not "dumb") agents
- iv. Episodic (vs. Sequential): agent's experience is divided into atomic "episodes", and choice of action in each episode depends only on the episode itself.
- v. Static (vs. dynamic): environment is unchanged while agent deliberates.
 - Semi-dynamic: if environment does not change with time, but agent's performance score does.
- vi. Discrete (vs. Continuous): a limited number of distinct percepts and actions.

Agent Structures

Agents are completely specified by the agent function:

- i. Agent Function: $f : \mathcal{P} \rightarrow \mathcal{A}$ maps from precept histories \mathcal{P} to actions \mathcal{A}
- ii. Agent Program: implements the agent function

Common Agent Structures:

- i. Simple Reflex: chooses action only based on current percept, ignoring precept history (follow if-then rules).
- ii. Goal-based: select actions to achieve a given tracked goal.
- iii. Utility-based: selects actions to maximise a utility function which assigns a score to any precept sequence.
If the utility function aligns with performance measure, the agent is rational.
- iv. Learning: improves performance over time with experience.
 - Performance Element: selects actions.
 - Learning Element: updates knowledge from feedback.
 - Critic: provides feedback on performance relative to a fixed performance standard.
 - Problem Generator: suggest exploratory actions

Exploration-Exploitation Dilemma:

- i. Explore: learn more about the world.
- ii. Exploit: maximize gain from current knowledge.

2. Systematic Search

Systematic search problems are fully-observable, deterministic, static, discrete and defined by:

- i. States: representation of problem state
- ii. Initial State
- iii. Goal State/Test
- iv. Actions: possible operations on a state
- v. Transition Model: function of action on a state
- vi. Action Cost Function

Uninformed Search

Uninformed search explores a problem space without domain-specific knowledge or heuristics to guide searching.

- i. Branching factor b : max successors of any node
- ii. Optimal solution depth d , Maximum depth m

Breadth-First Search (BFS) expands nodes level-by-level using a queue:

- i. Time: $O(b^{d+1})$
- ii. Space: $O(b^d)$
- iii. Complete: Yes (if b is finite)
- iv. Optimal: Yes (if all same step cost)

Uniform-Cost Search (UCS) expands least-cost node using a priority queue, for optimal solution cost C^* , levels C^*/ϵ :

- i. Time: $O(b^{C^*/\epsilon})$
- ii. Space: $O(b^{C^*/\epsilon})$
- iii. Complete: Yes (if step costs $> 0 \wedge$ finite total cost)
- iv. Optimal: Yes (if step costs > 0)

Depth-First Search (DFS) expands deepest unexpanded nodes using a stack:

- i. Time: $O(b^m)$
- ii. Space: $O(bm)$
- iii. Complete: No (if infinite depth \vee cyclic)
- iv. Optimal: No

Depth-Limited Search (DLS) limits search depth to ℓ :

- i. Time: $O(b^\ell)$
- ii. Space: $O(b\ell)$ (with DFS)
- iii. Complete: No (if $\ell < d$)
- iv. Optimal: No (with DFS)

Iterative Deepening Search (IDS) runs DLS with increasing depth limit until solution found:

- i. Time: $O(b^d)$ (with additional overhead)
- ii. Space: $O(bd)$ (with DFS)
- iii. Complete: Yes (if b finite)
- iv. Optimal: Yes (if all same step cost)

Informed Search

Informed search uses heuristics to guide searching, where heuristic $h(n)$ estimates optimal cost from a state to goal:

- i. All heuristics must be non-negative $\wedge h(goal) = 0$

Usefulness Properties:

- i. Admissible: \forall node n , $h(n) \leq h^*(n)$ where $h^*(n)$ is optimal path cost to reach goal state from n ($h(n)$ never over-estimates true cost to goal)
 - Theorem: if $h(n)$ is admissible, A^* search (without visited memory) is optimal
- ii. Consistent: \forall node n , $h(n) \leq c(n, a, n') + h(n')$ (triangle inequality)
 - Theorem: if $h(n)$ is consistent, A^* search (with visited memory) is optimal
 - Theorem: if $h(n)$ is consistent, $h(n)$ is admissible
- iii. Dominance: \forall nodes n , $h_1(n) \geq h_2(n)$ implies that h_1 dominates h_2 (more informed)
 - Theorem: if h_1 is admissible, h_1 is better for search, expanding no more nodes than h_2

Admissible heuristics can be found using optimal cost of a relaxed problem (fewer restrictions on actions):

- Ex. h_{SLD} straight-line distance if agent can fly
- Ex. $h = 0$ if agent can teleport

Best-First Search expands nodes with evaluation $f(n) = h(n)$ using a priority queue:

- i. Time: $O(b^m)$
- ii. Space: $O(b^m)$
- iii. Complete: No (may get stuck in loops).
- iv. Optimal: No (heuristic-only may mislead).

A^* expands nodes with evaluation $f(n) = g(n) + h(n)$, where $g(n)$ is step cost, using a priority queue:

- i. Time: $O(b^d)$ (good heuristic can improve)
- ii. Space: $O(b^d)$ (keeps all nodes in memory).
- iii. Complete: Yes (if step costs $> 0 \wedge$ finite b).
- iv. Optimal: Depends (if h admissible in tree search, or consistent in graph search).

3. Local Search

Local search problems are defined by:

- i. States: representation of candidate solution, may not map to actual problem state
- ii. Initial State
- iii. Goal State/Test (optional)
- iv. Successor Function: generate neighbour states by applying modifications to current state

Local search explores large, otherwise intractable problem spaces by considering only locally reachable states, guided by an evaluation function:

- i. Evaluation Function: assesses quality of a state; we either minimize or maximise this function

State Space Landscape:

- i. Global Maximum: optimal solution
- ii. Local Maximum: local optimum solution
- iii. Shoulder: region with flat evaluation function, difficult for algorithm to move past

Hill-Climbing searches for local optimum, generating successors from current state and picking the best using heuristic:

- i. Any-time: More time gives better solutions
- ii. Space: $O(b)$
- iii. Complete: No
- iv. Optimal: Not guaranteed
- v. Variants:
 - Simulated Annealing: allow some bad moves, gradually decreasing frequency
 - Beam Search: perform parallel k hill-climbs
 - Genetic Algorithm: successor generated by combining 2 parent states
 - Random-Restart: escape local optimum by restarting search

4. Adversarial Search

Adversarial search problems are fully-observable, deterministic-strategic, static, discrete and defined by:

- i. States: representation of candidate solution, may not map to actual problem state
- ii. Initial State
- iii. Terminal States: where game ends (e.g. win)
- iv. Actions: possible operations on a state
- v. Transition Model: function of action on a state
- vi. Utility Function: output value of state from the perspective of our agent

Minimax assumes both players play optimally, expanding game tree in DFS with alternating MAX and MIN levels:

- i. Time: $O(b^m)$
- ii. Space: $O(bm)$
- iii. Complete: Yes (for finite games)
- iv. Optimal: Yes (for both, if both play optimally)
 - Theorem: if opponent plays sub-optimally, utility obtained by agent is never less than utility obtained against an optimal opponent

Alpha-Beta Pruning reduces nodes evaluated without altering minimax value and optimal move for root node:

- i. Maintains α = best value, β = worst value and pruning subtree if $\alpha \geq \beta$
- ii. Time: $O(b^{m/2})$ (with perfect ordering)
- iii. Space: $O(bm)$ (same as minimax)
- iv. Complete & Optimal: same as minimax

Cutoff Strategy halts search in the middle and estimates value of mid-game states with an evaluation function:

- i. Evaluation Function: if terminal, use utility function, else use a heuristic
- ii. Heuristic Function: estimates utility of a state
 - A heuristic must be between max and min utility
 - Admissibility and consistency are not needed
- iii. Theorem: returns a move that is optimal with respect to the evaluation function at the cutoff; may be suboptimal with respect to true minimax

5. Machine Learning

Machine learning develops learning agents with data.

- i. Data, $D = \{1 \leq i \leq n : (x^{(i)}, y^{(i)})\}$ of n samples, for $x^{(i)} \in \mathbb{R}^d$ feature vector, $y^{(i)}$ label of i^{th} data point.
- ii. Model/Hypothesis, $h : X \rightarrow Y$ is a predictor of outputs which is learned by a learning algorithm.
- iii. Performance Measure evaluates h map $x^{(i)} \rightarrow y^{(i)}$.
 - $MAE = \frac{1}{n} \sum_{i=1}^n |\hat{y}^{(i)} - y^{(i)}|$ (outlier robust)
 - $Accuracy = \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{\hat{y}^{(i)} = y^{(i)}}$ (classification)
- iv. Loss Function measures $\hat{y}^{(i)}$ error for optimization
 - $MSE = \frac{1}{n} \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2$ (outlier sensitive)

Binary Metrics:

- i. Accuracy = $\frac{TP + TN}{TP + FN + FP + TN}$
- ii. Precision, $P = \frac{TP}{TP + FP}$ (if FP are costly)
- iii. Recall, $R = \frac{TP}{TP + FN}$ (if FN are costly)
- iv. F1 Score, $F1 = \frac{2}{\frac{1}{P} + \frac{1}{R}}$ (balance precision & recall)

Decision Trees

Decision trees split data using features to predict outputs.

- i. Representational Completeness: any decision tree can fit any finite consistent labelled data exactly

Entropy: $H(Y) = - \sum_i P(y_i) \log_2 P(y_i)$

Cond. Entropy: $H(Y | X) = \sum_x P(X=x)H(Y | X=x)$

Information Gain: $IG(Y; X) = H(Y) - H(Y | X)$

DTL grows tree top-down, greedily choosing feature X with highest IG, and recurses. At leaves:

- i. If no more data, return default
- ii. If data has same classification, return classification
- iii. If no more features, return majority class

Pruning reduces overfitting, giving simpler hypothesis by limiting representational capacity, removing outliers:

- i. Max-depth: limit max depth of decision tree.
- ii. Min-sample leaves: set min samples for leaf nodes.

6. Supervised Learning

Supervised Learning learns from labelled data to learn a mapping from inputs to outputs.

Linear Regression

Linear regression creates a linear model to predict $y \in \mathbb{R}$:

$$h_w(x) = w^T x = w_0 + w_1 x_1 + \dots + w_d x_d$$

where w_0, \dots, w_d are weights, using MSE loss function:

$$J_{MSE}(w) = \frac{1}{2n} \sum_{i=1}^n (h_w(x^{(i)}) - y^{(i)})^2$$

Learning uses $\frac{\partial J(w)}{\partial w_j} = \frac{1}{n} \sum_{i=1}^n (h_w(x^{(i)}) - y^{(i)}) x_j^{(i)}$ via:

- i. Normal Equation: $w = (X^\top X)^{-1} X^\top Y$
 - Closed form, assumes invertible, $O(d^3)$ inversion
- ii. Gradient Descent: $w_j \leftarrow w_j - \gamma \frac{\partial J(w)}{\partial w_j}$ repeated
 - Converge on global min., for appropriate $\gamma > 0$
 - If features linearly indep., global min is unique
 - May need feature scaling/ different γ_j for weights
 - Stochastic/ Mini-batch faster by using less data

Logistic Regression

Logistic regression creates a model to classify $y \in \{0, 1\}$:

$$h_w(x) = \sigma(w^T x), \quad \sigma(z) = \frac{1}{1 + e^{-z}}$$

where w is the weight vector and $\sigma(z)$ is sigmoid function with $\sigma'(z) = \sigma(z)(1 - \sigma(z))$, using BCE loss function:

$$J_{BCE} = \frac{1}{n} \sum_{i=1}^n BCE(y^{(i)}, h_w(x^{(i)}))$$

$$BCE(y, p) = -y \log(p) - (1 - y) \log(1 - p)$$

Model $h_w(x)$ can be interpreted as $P(y=1 | x)$, where we classify $y=1$ if $h_w(x) \geq \tau$, decision boundary $h_w(x) = \tau$

Learning only via Gradient Descent with same properties.

Feature Transformation $x \in \mathbb{R}^d \rightarrow \phi(x) \in \mathbb{R}^M$ allows linear and logistic regression on non-linear relationships.

Regularisation

Regularisation augments constraints to prevent overfitting:

- $J_{reg}(w) = J(w) + \lambda R(w)$, where $\lambda > 0$
- i. Generally, L_p -norm, $R_{L_p} = \|w\|_p = (\sum_{j=0}^d |w_j|^p)^{\frac{1}{p}}$
- ii. L_1 -norm, $R_{L_1} = \|w\|_1 = \sum_{j=0}^d |w_j|$
 - $\frac{\partial J_{L_1}^{MSE}(w)}{\partial w_j} = \frac{1}{n} \sum_{i=1}^n (w^\top x^{(i)} - y^{(i)}) x_j^{(i)} + \lambda \frac{\partial |w_j|}{\partial w_j}$
 - Theorem: features with non-zero weight will have constant penalty of λ
 - Theorem: L_1 induces sparsity, selecting features with $J_{MSE}(w) = \lambda$
- iii. L_2 -norm, $R_{L_2} = \|w\|_2 = \sqrt{\sum_{j=0}^d w_j^2}$ or $\frac{1}{2} \sum_{j=0}^d w_j^2$
 - $\frac{\partial J_{L_2}^{MSE}(w)}{\partial w_j} = \frac{1}{n} \sum_{i=1}^n (w^\top x^{(i)} - y^{(i)}) x_j^{(i)} + \lambda w_j$
 - Theorem: $X^\top X + \lambda I$ is invertible with unique, closed-form $w = (X^\top X + \lambda I)^{-1} X^\top Y$
 - Theorem: L_2 induces weights shrinkage
- iv. L_∞ -norm, $R_{L_\infty} = \|w\|_\infty = \max_j |w_j|$

Dual Formulation and Kernel Method

Dual formulation re-parameterises w using training data. Representer Theorem for optimal w of $J_{L_2}^{MSE}(w)$:

$$w = \sum_{j=1}^n \alpha_j x^{(j)}, \quad h_\alpha(x) = \sum_{j=1}^n \alpha_j x^{(j)\top} x = \sum_{j=1}^n \alpha_j \langle x^{(j)}, x \rangle$$

$$J_{L_2}^{MSE}(\alpha) = \frac{1}{2} \sum_{i=1}^n \left(\sum_{j=1}^n \alpha_j x^{(j)\top} x^{(i)} - y^{(i)} \right)^2 + \lambda \sum_{i=0}^n \left(\sum_{j=1}^n \alpha_j x_i^{(j)} \right)^2$$

Kernel function $k_\phi(u, v) = \langle \phi(u), \phi(v) \rangle$ exists for any feature mapping ϕ , giving cheaper computation:

$$h_\alpha^\phi(x) = \sum_{j=1}^n \alpha_j k_\phi(x^{(j)}, x)$$

Common kernels:

- i. Polynomial deg. k : $k_{pk}(u, v) = (u^\top v)^k$
- ii. Gaussian var. s^2 : $k_{RBF}(u, v) = \exp(-\frac{\|u-v\|^2}{2s^2})$

Support Vector Machines (SVMs)

SVMs are models to classify $y \in \{-1, 1\}$:

$$h_w(x) = \text{sign}(w^T x), \quad h_\alpha(x) = \text{sign}(\sum_{x^{(i)} \in S} \alpha_i k(x^{(i)}, x))$$

where w is normal vector of the zero-offset hyperplane decision boundary $w^\top x_{\in S} = 0$, maximising objective:

$$\max_w \frac{1}{\|w\|} = \min_w \frac{\|w\|^2}{2}, \quad y^{(i)}(w^\top x^{(i)}) \geq 1$$

$$\max_\alpha \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{i=j=1}^n \alpha_i \alpha_j y^{(i)} y^{(j)} k_\phi(x^{(i)}, x^{(j)}), \sum_{i=1}^n \alpha_i y^{(i)} = 0$$

- i. Distance to point $x: \frac{|w^\top x|}{\|w\|}$, Margin: $\frac{1}{\|w\|}$
- ii. Support vectors S lie exactly on margin boundary
 - $\alpha_i > 0 \iff x^{(i)} \in S, \alpha_i = 0 \iff x^{(i)} \notin S$
 - Theorem: for linearly separable data of $r \leq d$ effective dimension, $|S| \geq r + 1$
(i.e. support vectors much less than data points)

Applications

Multi-class Classification ($y \in \{1, \dots, K\}$):

- i. One-One: $\binom{K}{2}$ classifiers vote between I, J
- ii. One-Rest: K classifiers choose between I, I^c

Multi-label Classification ($y \in \{0, 1\}^K$):

- i. Binary Relevance: indep. binary classifiers
 - ignores label correlations and constraints
- ii. Classifier Chains: chain pred. of binary classifiers
 - captures dependencies; order-sensitive
- iii. Label Powerset
 - preserve correlations; classes explode; data-sparse

Generalizability:

- i. Dataset: relevance, noise, balance (classification)
- ii. Model Complexity:
 - Low: underfits for complex data, high bias, low variance on re-trains on different data
 - High: overfits for scarce data, low bias with enough data, high variance on re-trains on different data

7. Unsupervised Learning

Unsupervised Learning finds patterns in unlabelled data.

K-Means Clustering

K-means Clustering groups n data points into K groups with centroids μ_1, \dots, μ_K , where centroid μ_j of n_j points is the average of points $x^{(i)}$ in the cluster:

$$\mu_j = \frac{1}{n_j} \sum_{i=1}^{n_j} x^{(i)}$$

with evaluation function:

$$J(c^{(1)}, \dots, c^{(n)}, \mu_1, \dots, \mu_K) = \frac{1}{n} \sum_{i=1}^n \|x^{(i)} - \mu_{c^{(i)}}\|^2$$

1. Randomly initialize K centroids μ_1, \dots, μ_K
2. Repeat until convergence:
 - 2.1. For $i = 1, \dots, n$: $c^{(i)} \leftarrow k$ of closest μ_k
 - 2.2. For $k = 1, \dots, K$: $\mu_k \leftarrow$ centroid of data points $x^{(i)}$ assigned to cluster k
- i. Theorem: each iteration never increases distortion; need deterministic tie-break for conv. to local opt.
- ii. Choose K by business need or using elbow method (stop after decrease in J with K slows suddenly)
- iii. Can be used for classification.

Principal Component Analysis (PCA)

PCA preserves variance while reducing features to $r < d$:

- i. SVD of $d \times n X^\top$: $X^\top = U \Sigma V^\top$
 - U is $d \times d$ orthonormal columns and rows
 - Σ is $d \times n$ diagonal with ordered $\sigma_j \geq 0$
 - V is $n \times n$ orthonormal columns and rows
- ii. Truncating to σ_r yields $d \times r \tilde{U}, r \times n \tilde{\Sigma}$
- iii. Compression of X^\top : $r \times n Z = \tilde{U}^\top X^\top$
- iv. Reconstruction: $\hat{X}^\top \approx \tilde{U} \tilde{U}^\top X^\top$ since $\tilde{U} \tilde{U}^\top \approx I$ with approximation error dependent on r
- v. Theorem: in mean-centered \hat{X}^\top , $\frac{\sigma_j^2}{n-1}$ are variance of data in the basis of $u^{(j)}$
- vi. Var: $\frac{\sum_{i=1}^r \sigma_i^2}{\sum_{i=1}^n \sigma_i^2} \geq 0.99 \iff \frac{\sum_{i=1}^n \|\hat{x}^{(i)} - \tilde{x}^{(i)}\|^2}{\sum_{i=1}^n \|\hat{x}^{(i)}\|^2} \leq 0.01$

8. Neural Networks

Neural networks are layered models of neurons that learn feature transformations for prediction tasks.

Neuron computes activation function on a weighted sum:

$$\hat{y} = g(z), \quad z = \sum_{j=0}^d w_j x_j = w^\top x$$

- i. Common activation functions:

- Identity: $g(z) = z$
- Sigmoid: $g(z) = \sigma(z) = \frac{1}{1+e^{-z}}$
- Tanh: $g(z) = \tanh(z) = 2\sigma(2x) - 1 = \frac{e^z - e^{-z}}{e^z + e^{-z}}$
- Relu: $g(z) = \max(0, z)$
- Leaky Relu: $g(z) = \max(az, z)$
- Softmax: $g(z_i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}} \in [0, 1]$

- ii. Output Layer: chosen to match task target
- iii. Hidden Layer: any layer between input and output

Forward Propagation is the process of input data passing through neural network to output layer:

- i. Weights of layer i , $W^{[i]}$, where $W_{jk}^{[i]}$ are weights from previous neuron j to layer neuron k
- ii. Activation function of layer i , $g^{[i]}$
- iii. Output of layer i , $\hat{y}^{[i]} = g^{[i]}(W^{[i]\top} y^{[i-1]})$

Backpropagation is gradient descent in the network:

- i. Compute loss $L(\hat{y}, y)$ at output layer
- ii. Apply chain rule layer-by-layer backwards for $\frac{\partial L}{\partial w_j} = \frac{dL}{dy} \frac{\partial y}{\partial z} \frac{\partial z}{\partial w_j}$ and $\frac{dL}{dx} = \frac{\partial L}{\partial z_1} \frac{\partial z_1}{\partial x} + \frac{\partial L}{\partial z_2} \frac{\partial z_2}{\partial x}$
- iii. Algorithm: Compute $\frac{dL}{d\hat{y}}$, convert $g(z)$ to $g'(z)$, reverse graph and compute gradients

Tasks:

- i. Binary Classification: 1 sigmoid output neuron
- ii. Multi-class Classification: K softmax output neurons for $\hat{y} = [\hat{y}_1, \dots, \hat{y}_K]^\top$ as one-hot vector
- iii. Single Regression: 1 any output neuron
- iv. Multi Regression: K any output neurons for $\hat{y} = [\hat{y}_1, \dots, \hat{y}_K]^\top$

Convolutional Neural Networks (CNNs)

CNNs apply shared filters over local regions of inputs:

- Convolution Layer: slide kernel over input and compute weighted sums to produce feature maps
 - Kernel/Filter: matrix of weights
 - Stride: step size
 - Padding: add k padding around border of input
 - Output: $N' = \lfloor \frac{N-K+2P}{S} \rfloor + 1$
- Pooling Layer: downsample feature maps to reduce resolution and parameters
 - Max-Pool: keep max value in window
 - Average-Pool: compute average value in window
- Architecture:
 - Feature Transformation: stack of Conv and Pools
 - Prediction: fully-connected output layer

usually using CE loss between predicted \hat{y} and actual y :

$$J_{CE} = \frac{1}{n} \sum_{i=1}^n CE(y^{(i)}, \hat{y}^{(i)})$$

$$CE(y^{(i)}, \hat{y}^{(i)}) = - \sum_{j=1}^K y_j^{(i)} \log(\hat{y}_j^{(i)})$$

Recurrent Neural Networks (RNNs)

RNNs model sequences using hidden state, where for layer j with hidden state $h^{[t]}$ at time t :

$$h^{[t]} = \hat{y}^{[j]} = g^{[j]}(W^{[ij]\top} x^{[t]} + W^{[hj]\top} h^{[t-1]})$$

Tasks:

- Many-Many (e.g. sequence tagging, translation):
1 input and 1 output neuron for $t \leq T_x = T_y$
- Many-Many (e.g. QnA):
1 input for $t \leq T_x$, "BEGIN" at $t = T_x + 1$, and 1 output neuron for $T_x < t \leq T_y$
- Many-One (e.g. sentiment analysis):
1 input for $t \leq T_x$, 1 output at final $t = T_y = T_x$
- One-Many (e.g. image captioning):
1 input at start $t = T_x = 1$, 1 output for $t \leq T_y$

Attention Neural Networks (ANNs)

Attention lets a model selectively focus on relevant parts of a sequence using vectors derived from $d \times T$ input X :

- Query, $d_q \times T Q$: represent current element query
- Key, $d_k \times T K$: contain keys to search information
- Value, $d_v \times T V$: actual element information

Attention Layers:

- Self-Attention: Q, K, V from same sequence
 - $q^{[t]} = W^q x^{[t]} \iff Q = W^q X$, W^q is $d_q \times d$
 - $k^{[t]} = W^k x^{[t]} \iff K = W^k X$, W^k is $d_k \times d$
 - $v^{[t]} = W_v x^{[t]} \iff V = W^v X$, W^v is $d_v \times d$
 - $\alpha_{it} = \frac{(k^{[i]})^\top q^{[t]}}{\sqrt{d_k}} \iff A = \frac{K^\top Q}{\sqrt{d_k}}$, A is $T \times T$
 - $\alpha'_{it} = \text{softmax}(\alpha_{it}) \iff A' = \text{softmax}(A)$ col-w
 - $h^{[t]} = \sum_{i=1}^T \alpha'_{it} v^{[i]} \iff H = VA'$, H is $d_v \times T$
- Masked Self-Attention: restrict to only $\leq t$
 - $a_{it} = \begin{cases} a_{it}, & \text{if } i \leq t \\ -\infty, & \text{if } i > t \end{cases}$
- Cross-Attention: Q from decoder sequence, K, V from encoder sequence
 - $k^{[t]} = W^k z^{[t]}$ i.e. $K = W^k Z$
 - $v^{[t]} = W_v z^{[t]}$ i.e. $V = W^v Z$

Tasks:

- Many-Many: self-attention layer with BCE loss
- Positional Encoding (distinct, consistent with T , bounded): $x'^{[t]} = x^{[t]} + PE^{[t]}$ where for pos. k :

$$PE_k^{[t]} = \begin{cases} \sin(t/C^{\frac{k}{d}}), & \text{if } k \% 2 == 0 \\ \cos(t/C^{\frac{k-1}{d}}), & \text{if } k \% 2 == 1 \end{cases}$$
- Many-Many (e.g. QnA): self-attention for question, masked self-attention and cross-attention for answer
- Many-One (e.g. sentiment): self-attention layer with additional "CLS" input for $h^{[CLS]}$ to capture summary information of input sequence
- One-Many (e.g. image caption): masked self-attention layer with additional "BEGIN" input

Teacher Forcing is used in training for sequence prediction:

- Input for next step is true target, not models own prediction from current step
- All output can be generated simultaneously, no need to wait for previous word
- In early training, predictions are nearly random, which may make the model lost
- During testing, teacher forcing is not available, leading to train-test mismatch
- Mitigation: sometimes feed true word, sometimes prediction, gradually reducing teacher forcing

Transformers

Transformers are deep attention neural networks:

- Stack of Encoder Blocks; self-attention \rightarrow feed-forward
- Stack of Decoder Blocks; masked self-attention \rightarrow cross-attention \rightarrow feed-forward

Unsupervised Learning with Neural Networks

Autoencoders use neural networks for dimensionality reduction and representation learning:

- Encoder: converts high-dimensional d input to $k < d$ dimensional compressed data.
- Decoder: reconstructs original data from compressed data.
- Reconstruction Loss: measures difference between reconstructed output and original input.
- After Training: discard decoder, use encoder to generate compressed representation for each input.

Model Pre-training initialises feature transformation weights based on tasks which are solvable with unlabelled data, challenging enough to force understanding of task, and yield useful learned understanding:

- Image Rotation Prediction
- Contrastive Learning (positive, negative pairs)
- Image Inpainting
- Next-word Prediction

Problems with Neural Networks

Common problems when training deep neural networks:

i. Overfitting:

- Model fits training data too well and performs poorly on unseen data.
- Early Stopping: stop training when performance on validation set begins to worsen.
- Dropout: during training, randomly set some neurons' output to 0; prevents overfitting by making the network less reliant on specific neurons.

ii. Vanishing / Exploding Gradient:

- Vanishing gradient: gradients become very small, causing very slow learning in deep networks.
- Exploding gradient: gradients become very large, causing unstable updates.
- Gradient Clipping: clip gradients within range $[-\text{clip_value}, \text{clip_value}]$ to control exploding gradients.

Perceptrons

Perceptron is a single neuron for binary classification:

$$\hat{y} = \text{sign}(w^\top x + b)$$

- Represents a linear decision boundary (hyperplane)
- Perceptron update rule for misclassified i :
 $w' \leftarrow w + \lambda(y^{(i)} - \hat{y}^{(i)})x^{(i)}$
- Multi-Layer Perceptron (MLP): stack of fully-connected layers with non-linear activations.

Receptive Fields in CNNs

Receptive field of a neuron is the region of input that can affect its activation.

Formualae for receptive field r_i for a neuron in layer i :

$$r_i = r_{i-1} + (K_i - 1) \times j_{i-1}$$

$$j_i = j_{i-1} \times S_i$$

where K_i is kernel size at i , S_i is stride at i , $r_0 = j_0 = 1$

```

def alpha_beta_search(state):
    v = max_value(state, -∞, ∞)
    return action in successors(state) with value v

def max_value(state, α, β):
    if is_terminal(state): return utility(state)
    v = -∞
    for next_state in expand(state):
        v = max(v, min_value(next_state, α, β))
        α = max(α, v)
        if v >= β: return v
    return v

def min_value(state, α, β):
    if is_terminal(state): return utility(state)
    v = ∞
    for next_state in expand(state):
        v = min(v, max_value(next_state, α, β))
        β = min(β, v)
        if v <= α: return v
    return v

```

a-B pruning Steps:

- Start at root $\alpha = -\infty, \beta = \infty$
- When going down, pass down the values of α, β
- When going up, pass up the value of α if MAX, or β if MIN
- When taking a value up, store $\alpha = \max(\alpha, \text{value})$ if MAX, or $\beta = \min(\beta, \text{value})$ if MIN
- At any point, if there is $\alpha \geq \beta$ prune all other children

Worst-case

Name	Time Complexity ¹	Space Complexity ¹	Complete?	Optimal?
Breadth-first Search	Exponential	Exponential	Yes	Yes
Uniform-cost Search	Exponential	Exponential	Yes ³	Yes ³
Depth-first Search	Exponential	Polynomial	No Yes ⁴	No
Depth-limited Search	Exponential	Polynomial ²	No	No ²
Iterative Deepening Search	Exponential	Polynomial ²	Yes	Yes
A* Search	Exponential	Exponential	Yes ³	Yes if admissible
A* Search with visited memory	Exponential	Exponential	Yes ³	Yes if consistent

1) In terms of some notion of depth/tier

2) If used with DFS

3) If edge costs are positive

4) With visited memory