


## 1. Intelligent Agents

Agents receive precepts from the environment through sensors and perform actions through actuators.

### PEAS Framework

PEAS defines problems in terms of Performance Measure, Environment, Actuators and Sensors:

- i. Rational Agent: chooses actions that maximise performance measure



Properties of Task Environment:

- i. Fully Observable (vs. Partially Observable): sensors give access to the complete state of the environment at each point in time.
- ii. Single-agent (vs. Multi-agent): agent operates by itself in an environment.
- iii. Deterministic (vs. Stochastic): next state of the environment is completely determined by current state and action by agent.
  - Strategic: if environment is dependent on actions of other unpredictable (not "dumb") agents
- iv. Episodic (vs. Sequential): agent's experience is divided into atomic "episodes", and choice of action in each episode depends only on the episode itself.
- v. Static (vs. dynamic): environment is unchanged while agent deliberates.
  - Semi-dynamic: if environment does not change with time, but agent's performance score does.
- vi. Discrete (vs. Continuous): a limited number of distinct percepts and actions.

## Agent Structures

Agents are completely specified by the agent function:

- i. Agent Function:  $f : \mathcal{P} \rightarrow \mathcal{A}$  maps from precept histories  $\mathcal{P}$  to actions  $\mathcal{A}$
- ii. Agent Program: implements the agent function

Common Agent Structures:

- i. Simple Reflex: chooses action only based on current percept, ignoring precept history (follow if-then rules).
- ii. Goal-based: select actions to achieve a given tracked goal.
- iii. Utility-based: selects actions to maximise a utility function which assigns a score to any precept sequence. If the utility function aligns with performance measure, the agent is rational.
- iv. Learning: improves performance over time with experience.
  - Performance Element: selects actions.
  - Learning Element: updates knowledge from feedback.
  - Critic: provides feedback on performance relative to a fixed performance standard.
  - Problem Generator: suggest exploratory actions

Exploration-Exploitation Dilemma:

- i. Explore: learn more about the world.
- ii. Exploit: maximize gain from current knowledge.

## 2. Systematic Search

Systematic search problems are fully-observable, deterministic, static, discrete and defined by:

- i. States: representation of problem state
- ii. Initial State
- iii. Goal State/Test
- iv. Actions: possible operations on a state
- v. Transition Model: function of action on a state
- vi. Action Cost Function

## Uninformed Search

Uninformed search explores a problem space without domain-specific knowledge or heuristics to guide searching.

- i. Branching factor  $b$ : max successors of any node
- ii. Optimal solution depth  $d$ , Maximum depth  $m$

Breadth-First Search (BFS) expands nodes level-by-level using a queue:

- i. Time:  $O(b^{d+1})$
- ii. Space:  $O(b^d)$
- iii. Complete: Yes (if  $b$  is finite)
- iv. Optimal: Yes (if all same step cost)

Uniform-Cost Search (UCS) expands least-cost node using a priority queue, for optimal solution cost  $C^*$ , levels  $C^*/\epsilon$ :

- i. Time:  $O(b^{C^*/\epsilon})$
- ii. Space:  $O(b^{C^*/\epsilon})$
- iii. Complete: Yes (if step costs  $> 0 \wedge$  finite total cost)
- iv. Optimal: Yes (if step costs  $> 0$ )

Depth-First Search (DFS) expands deepest unexpanded nodes using a stack:

- i. Time:  $O(b^m)$
- ii. Space:  $O(bm)$
- iii. Complete: No (if infinite depth  $\vee$  cyclic)
- iv. Optimal: No

Depth-Limited Search (DLS) limits search depth to  $\ell$ :

- i. Time:  $O(b^\ell)$
- ii. Space:  $O(b\ell)$  (with DFS)
- iii. Complete: No (if  $\ell < d$ )
- iv. Optimal: No (with DFS)

Iterative Deepening Search (IDS) runs DLS with increasing depth limit until solution found:

- i. Time:  $O(b^d)$  (with additional overhead)
- ii. Space:  $O(bd)$  (with DFS)
- iii. Complete: Yes (if  $b$  finite)
- iv. Optimal: Yes (if all same step cost)

## Informed Search

Informed search uses heuristics to guide searching, where heuristic  $h(n)$  estimates optimal cost from a state to goal:

- i. All heuristics must be non-negative  $\wedge h(goal) = 0$

Usefulness Properties:

- i. Admissible:  $\forall$  node  $n$ ,  $h(n) \leq h^*(n)$  where  $h^*(n)$  is optimal path cost to reach goal state from  $n$  ( $h(n)$  never over-estimates true cost to goal)
  - Theorem: if  $h(n)$  is admissible,  $A^*$  search (without visited memory) is optimal
- ii. Consistent:  $\forall$  node  $n$ ,  $h(n) \leq c(n, a, n') + h(n')$  (triangle inequality)
  - Theorem: if  $h(n)$  is consistent,  $A^*$  search (with visited memory) is optimal
  - Theorem: if  $h(n)$  is consistent,  $h(n)$  is admissible
- iii. Dominance:  $\forall$  nodes  $n$ ,  $h_1(n) \geq h_2(n)$  implies that  $h_1$  dominates  $h_2$  (more informed)
  - Theorem: if  $h_1$  is admissible,  $h_1$  is better for search, expanding no more nodes than  $h_2$

Admissible heuristics can be found using optimal cost of a relaxed problem (fewer restrictions on actions):

Ex.  $h_{SLD}$  straight-line distance if agent can fly

Ex.  $h = 0$  if agent can teleport

Best-First Search expands nodes with evaluation  $f(n) = h(n)$  using a priority queue:

- i. Time:  $O(b^m)$
- ii. Space:  $O(b^m)$
- iii. Complete: No (may get stuck in loops).
- iv. Optimal: No (heuristic-only may mislead).

$A^*$  expands nodes with evaluation  $f(n) = g(n) + h(n)$ , where  $g(n)$  is step cost, using a priority queue:

- i. Time:  $O(b^d)$  (good heuristic can improve)
- ii. Space:  $O(b^d)$  (keeps all nodes in memory).
- iii. Complete: Yes (if step costs  $> 0 \wedge$  finite  $b$ ).
- iv. Optimal: Depends (if  $h$  admissible in tree search, or consistent in graph search).

## 3. Local Search

Local search problems are defined by:

- i. States: representation of candidate solution, may not map to actual problem state
- ii. Initial State
- iii. Goal State/Test (optional)
- iv. Successor Function: generate neighbour states by applying modifications to current state

Local search explores large, otherwise intractable problem spaces by considering only locally reachable states, guided by an evaluation function:

- i. Evaluation Function: assesses quality of a state; we either minimize or maximise this function

State Space Landscape:

- i. Global Maximum: optimal solution
- ii. Local Maximum: local optimum solution
- iii. Shoulder: region with flat evaluation function, difficult for algorithm to move past

Hill-Climbing searches for local optimum, generating successors from current state and picking the best using heuristic:

- i. Any-time: More time gives better solutions
- ii. Space:  $O(b)$
- iii. Complete: No
- iv. Optimal: Not guaranteed
- v. Variants:
  - Simulated Annealing: allow some bad moves, gradually decreasing frequency
  - Beam Search: perform parallel  $k$  hill-climbs
  - Genetic Algorithm: successor generated by combining 2 parent states
  - Random-Restart: escape local optimum by restarting search

## 4. Adversarial Search

Adversarial search problems are fully-observable, deterministic-strategic, static, discrete and defined by:

- i. States: representation of candidate solution, may not map to actual problem state
- ii. Initial State
- iii. Terminal States: where game ends (e.g. win)
- iv. Actions: possible operations on a state
- v. Transition Model: function of action on a state
- vi. Utility Function: output value of state from the perspective of our agent

Minimax assumes both players play optimally, expanding game tree in DFS with alternating MAX and MIN levels:

- i. Time:  $O(b^m)$
- ii. Space:  $O(bm)$
- iii. Complete: Yes (for finite games)
- iv. Optimal: Yes (for both, if both play optimally)
  - Theorem: if opponent plays sub-optimally, utility obtained by agent is never less than utility obtained against an optimal opponent

Alpha-Beta Pruning reduces nodes evaluated without altering minimax value and optimal move for root node:

- i. Maintains  $\alpha$  = best value,  $\beta$  = worst value and pruning subtree if  $\alpha \geq \beta$
- ii. Time:  $O(b^{m/2})$  (with perfect ordering)
- iii. Space:  $O(bm)$  (same as minimax)
- iv. Complete & Optimal: same as minimax

Cutoff Strategy halts search in the middle and estimates value of mid-game states with an evaluation function:

- i. Evaluation Function: if terminal, use utility function, else use a heuristic
- ii. Heuristic Function: estimates utility of a state
  - A heuristic must be between max and min utility
  - Admissibility and consistency are not needed
- iii. Theorem: returns a move that is optimal with respect to the evaluation function at the cutoff; may be suboptimal with respect to true minimax

## 5. Machine Learning

Machine learning develops learning agents with data.

- i. Data,  $D = \{1 \leq i \leq n : (x^{(i)}, y^{(i)})\}$  of  $n$  samples, for  $x^{(i)} \in \mathbb{R}^d$  feature vector,  $y^{(i)}$  label of  $i^{th}$  data point.
- ii. Model/Hypothesis,  $h : X \rightarrow Y$  is a predictor of outputs which is learned by a learning algorithm.
- iii. Performance Measure evaluates  $h$  map  $x^{(i)} \rightarrow y^{(i)}$ .
  - $MAE = \frac{1}{n} \sum_{i=1}^n |\hat{y}^{(i)} - y^{(i)}|$  (outlier robust)
  - $Accuracy = \frac{1}{n} \sum_{i=1}^n \mathbb{1}_{\hat{y}^{(i)}=y^{(i)}}$  (classification)
- iv. Loss Function measures  $\hat{y}^{(i)}$  error for optimization
  - $MSE = \frac{1}{n} \sum_{i=1}^n (\hat{y}^{(i)} - y^{(i)})^2$  (outlier sensitive)

Binary Metrics:

- i. Accuracy =  $\frac{TP + TN}{TP + FN + FP + TN}$
- ii. Precision,  $P = \frac{TP}{TP + FP}$  (if FP are costly)
- iii. Recall,  $R = \frac{TP}{TP + FN}$  (if FN are costly)
- iv. F1 Score,  $F1 = \frac{2}{\frac{1}{P} + \frac{1}{R}}$  (balance precision & recall)

## Decision Trees

Decision trees split data using features to predict outputs.

- i. Representational Completeness: any decision tree can fit any finite consistent labelled data exactly

Entropy:  $H(Y) = - \sum_i P(y_i) \log_2 P(y_i)$

Cond. Entropy:  $H(Y | X) = \sum_x P(X = x) H(Y | X = x)$

Information Gain:  $IG(Y; X) = H(Y) - H(Y | X)$

DTL grows tree top-down, greedily choosing feature  $X$  with highest IG, and recurses. At leaves:

- i. If no more data, return default
- ii. If data has same classification, return classification
- iii. If no more features, return majority class

Pruning reduces overfitting, giving simpler hypothesis by limiting representational capacity, removing outliers:

- i. Max-depth: limit max depth of decision tree.
- ii. Min-sample leaves: set min samples for leaf nodes.

## 6. Supervised Learning

Supervised Learning learns from labelled data to learn a mapping from inputs to outputs.

### Linear Regression

Linear regression creates a linear model to predict  $y \in \mathbb{R}$ :

$$h_w(x) = w^T x = w_0 + w_1 x_1 + \dots + w_d x_d$$

where  $w_0, \dots, w_d$  are weights, using MSE loss function:

$$J_{MSE}(w) = \frac{1}{2n} \sum_{i=1}^n (h_w(x^{(i)}) - y^{(i)})^2$$

Learning uses  $\frac{\partial J(w)}{\partial w_j} = \frac{1}{n} \sum_{i=1}^n (h_w(x^{(i)}) - y^{(i)}) x_j^{(i)}$  via:

- i. Normal Equation:  $w = (X^T X)^{-1} X^T Y$ 
  - Closed form, assumes invertible,  $O(d^3)$  inversion
- ii. Gradient Descent:  $w_j \leftarrow w_j - \gamma \frac{\partial J(w)}{\partial w_j}$  repeated
  - Converge on global min., for appropriate  $\gamma > 0$
  - If features linearly indep., global min is unique
  - May need feature scaling/ different  $\gamma_j$  for weights
  - Stochastic/ Mini-batch faster by using less data

### Logistic Regression

Logistic regression creates a model to classify  $y \in \{0, 1\}$ :

$$h_w(x) = \sigma(w^T x), \quad \sigma(z) = \frac{1}{1 + e^{-z}}$$

where  $w$  is the weight vector and  $\sigma(z)$  is sigmoid function with  $\sigma'(z) = \sigma(z)(1 - \sigma(z))$ , using BCE loss function:

$$J_{BCE} = \frac{1}{n} \sum_{i=1}^n BCE(y^{(i)}, h_w(x^{(i)}))$$

$$BCE(y, p) = -y \log(p) - (1 - y) \log(1 - p)$$

Model  $h_w(x)$  can be interpreted as  $P(y = 1 | x)$ , where we classify  $y = 1$  if  $h_w(x) \geq \tau$ , decision boundary  $h_w(x) = \tau$

Learning only via Gradient Descent with same properties.

Feature Transformation  $x \in \mathbb{R}^d \rightarrow \phi(x) \in \mathbb{R}^M$  allows linear and logistic regression on non-linear relationships.

### Multi-class Classification ( $y \in \{1, \dots, K\}$ )

One-vs-One classifies every pair of classes:

- i.  $\binom{K}{2}$  classifiers choosing between class  $I$  and  $J$
- ii. Each classifier votes for a class, most votes wins

One-vs-Rest classifies per class vs. all others:

- i.  $K$  classifiers choosing between class  $I$  and not- $I$
- ii. Classifier with highest probability determines class

### Multi-label Classification ( $y \in \{0, 1\}^K$ )

Binary Relevance uses independent label binary classifiers:

- i. Ignores label correlations and constraints (e.g. mutual exclusive)

Classifier Chains feeds predictions of previous binary labels classifiers as features for subsequent labels:

- i. Captures label dependencies; sensitive to chain order

Label Powerset treats each unique label-set as a single multi-class label:

- i. Preserves correlations; can explode in classes and be data-sparse

### Generalization

Generalization is the model's performance on unseen data.

Dataset Factors:

- i. Relevance: features should be relevant to problem
- ii. Noise: outliers can hinder learning, generalization
- iii. Balance (for classification): ensures all classes are fairly represented for generalization over classes

Model Complexity (optimal complexity minimises error):

- i. Low Complexity: underfits for complex data, high bias; low variance on retraining across different training sets
- ii. High Complexity: overfits for scarce data; low bias with enough data; high variance across retraining on different training sets

## Support Vector Machines

## 7. Neural Networks

## 8. Unsupervised Learning

Unsupervised Learning learns from unlabelled data to find patterns or structure.

### K-Means Clustering

### Hierarchical Clustering

### Dimensionality Reduction

## 9. Ethics

```
def alpha_beta_search(state):  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v
```

```
def max_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for next_state in expand(state):  
        v = max(v, min_value(next_state, α, β))  
        α = max(α, v)  
        if v >= β: return v  
    return v
```

```
def min_value(state, α, β):  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for next_state in expand(state):  
        v = min(v, max_value(next_state, α, β))  
        β = min(β, v)  
        if v <= α: return v  
    return v
```

a-B pruning Steps:

- Start at root  $\alpha = -\infty, \beta = \infty$
- When going down, pass down the values of  $\alpha, \beta$
- When going up, pass up the value of  $\alpha$  if MAX, or  $\beta$  if MIN
- When taking a value up, store  $\alpha = \max(\alpha, value)$  if MAX, or  $\beta = \min(\beta, value)$  if MIN
- At any point, if there is  $\alpha \geq \beta$  prune all other children

7

- 1) In terms of some notion of depth/tier
- 2) If used with DFS
- 3) If edge costs are positive
- 4) With visited memory

Worst-case				
Name	Time Complexity <sup>1</sup>	Space Complexity <sup>1</sup>	Complete?	Optimal?
Breadth-first Search	Exponential	Exponential	Yes	Yes
Uniform-cost Search	Exponential	Exponential	Yes <sup>3</sup>	Yes <sup>3</sup>
Depth-first Search	Exponential	Polynomial	No   Yes <sup>4</sup>	No
Depth-limited Search	Exponential	Polynomial <sup>2</sup>	No	No <sup>2</sup>
Iterative Deepening Search	Exponential	Polynomial <sup>2</sup>	Yes	Yes
A * Search	Exponential	Exponential	Yes <sup>3</sup>	Yes if admissible
A * Search with visited memory	Exponential	Exponential	Yes <sup>3</sup>	Yes if consistent