

## 1. Introduction

Parallel Computing is the simultaneous use of multiple processing units to solve problems efficiently.

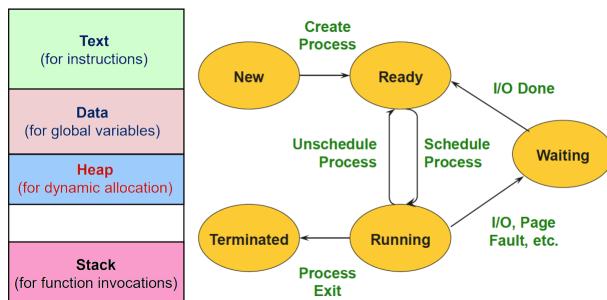
Program Parallelization Steps:

- Decomposition: split problem into parallel tasks
  - Granularity: size of task
- Scheduling: assign tasks to processes/threads
  - Orchestration: imposed synchronisation and communication of tasks to satisfy dependencies
- Mapping: bind processes/threads to hardware processing units (e.g. CPU cores)

## Processes

Process is an abstraction for a running program:

- Process ID (PID): uniquely identifies a process
- Process State: indicates execution status
- Process Control Block: stores execution context (registers, resources, exclusive address space)
- Costly: syscall overhead, data structures must be allocated, communication goes through OS



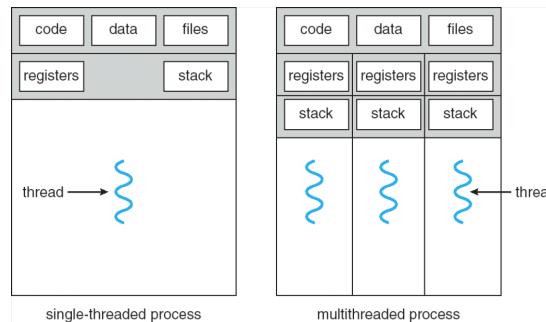
Interprocess Communication (IPC):

- Shared Memory (e.g. locks, semaphores)
- Message Passing
- UNIX Pipes and Signals

## Threads

Threads are independent execution flows within a process:

- Shared: Resources, Address Space
- Private: Thread ID, Registers, "Stack" (diff. SP)
- + Efficient: cheaper creation and context switching
- + Resource Sharing: process resources can be shared
- + Concurrent: multithreading on different cores



Implementations:

- User-Threads are managed by a userspace library:
  - + Fast: context switches have low overhead
  - No Parallelism: OS cannot map threads to different execution units
  - Threads blocking will block whole process
- Kernel-Threads are implemented in the OS:
  - + Parallel: threads can run across many CPUs
  - Slower: thread operations are system calls

Mapping Models:

- Many-to-One:
  - All user-level threads are mapped to one process
  - User library schedules user-threads
- One-to-One:
  - Each user-thread mapped to one kernel-thread
  - No user library scheduler needed
- Many-to-Many:
  - User-threads mapped to set of kernel-threads
  - Library scheduler may move user-threads to different kernel-threads

## Synchronisation

Synchronisation is required by concurrent threads to coordinate access to shared resources in critical sections.

Critical Section Properties:

- Safety: nothing bad happens
  - Mutual Exclusion: at most 1 thread in CS
- Liveness: something good happens
  - Progress: if no thread is in CS, a waiting thread should be granted access
  - Bounded Wait: a waiting thread requesting to enter the CS, will eventually enter
- Performance: overhead of entering/exiting CS is small w.r.t work done within it

Symptoms of Incorrect Synchronization:

- Deadlock: all threads blocked; iff all conditions met
  - Mutual Exclusion:  $\geq 1$  resource held exclusively
  - Hold & Wait:  $\geq 1$  process holding one resource while waiting for another
  - No Pre-emption: resources must be yielded
  - Circular Wait: set of processes waiting in circles
- Livelock: due to deadlock avoidance, no progress
- Starvation: a thread is unable to access CS
- Race Condition: bug from program outcome depending on unpredictable access order

Mechanisms:

- Lock: primitive `acquire()`, `release()`
  - Spinlock: busy-wait; wastes cycles
  - Mutex: blocking
  - Uses hardware atomics (TSL) or toggle interrupts
- Semaphore: atomic counter  $\geq 0$ ; `wait()`, `signal()`
  - Binary/Mutex: represents single access
  - Counting/General: number of threads in CS  $\leq N$
  - Can be manipulated by different thread
  - No relation to data being controlled
- Monitors: thread-safe high-level data structures
- Messages

## Producer-Consumer

Processes share a bounded buffer of fixed size  $K$ , where producers add items until buffer full and consumers remove items when not empty.

Blocking Solution (init `not_full = K, not_emp = 0`):

```
// Producer           // Consumer
x = produce();
wait(not_emp);
wait(mutex);
wait(not_full);
wait(mutex);
buffer.add(x);
signal(mutex);
signal(not_emp);

// Consumer
wait(not_full);
wait(mutex);
x = buffer.get();
signal(mutex);
signal(empty);
consume(item);
```

Busy Waiting Solution:

```
// Producer
while(!can_prod);
x = produce();

wait(mutex);
if (K > count) {
    buf[in] = x;
    in = (in+1)%K;
    count++;
    can_cons = 1;
} else {
    can_prod = 0;
}
signal(mutex);

// Consumer
while(!can_cons);

wait(mutex);
if (count > 0) {
    x = buf[out];
    out = (out+1)%K;
    count--;
    can_prod = 1;
} else {
    can_cons = 0;
}
signal(mutex);

consume(x);
```

## Reader-Writer

Processes share a critical region, where readers can read simultaneously but writers must have exclusive access.

```
// Reader
wait(mutex);
readers++;
if (readers == 1)
    wait(empty);
signal(mutex);

read();

wait(mutex);
readers--;
if (readers == 0)
    signal(empty);
signal(mutex);

// Writer
wait(empty);
write();
signal(empty);
```

- Writer Starvation: possible if readers keep entering and `empty` is never signalled. Solutions:
  - Queue: all pass through `wait(queue)`
  - Writer-Priority: readers entering must `wait(readTry)`; but first waiting writer blocks new readers with `wait(readTry)` and does `signal(readTry)` when done

## Barrier

```
int i_am_last = 0;
wait(mutex);
count++;
if (count == N) {
    i_am_last = 1;
    signal(barrier);
}
signal(mutex);

wait(barrier);
if (!i_am_last)
    signal(barrier);
```

## Dining Philosophers

$N$  philosophers around a circular table, with a single chopstick between. 2 chopsticks are needed to eat.

```
// Philosopher
void p(int i) {
    while (1) {
        // think
        take(i);
        eat();
        put(i);
    }
}

void safe_to_eat
    (int i) {
    if (state[i] == HUNGRY &&
        state[LEFT] != EATING &&
        state[RIGHT] != EATING) {
        state[i] = EATING;
        signal(s[i]);
    }
}
```

```
// Take Chopstick
void take(int i) {
    wait(mutex);
    state[i] = HUNGRY;
    safe_to_eat(i);
    signal(mutex);
    wait(s[i]);
}

// Put Chopstick
void put(int i) {
    wait(mutex);
    state[i] = THINKING;
    safe_to_eat(LEFT);
    safe_to_eat(RIGHT);
    signal(mutex);
}
```

Limited Eater Solution, when at most  $N - 1$  philosophers eat concurrently, it's guaranteed at least one can eat:

```
void philosopher(int i) {
    while (1) {
        // think
        wait(seats);
        wait(chopstick[LEFT]);
        wait(chopstick[RIGHT]);
        eat();
        signal(chopstick[RIGHT]);
        signal(chopstick[LEFT]);
        signal(seats);
    }
}
```

## 2. Architecture

Levels of Parallelism:

- Bit-Level: operate on wider words
- Instruction-Level:
  - Pipeline: split execution into stages; multiple instructions can occupy different stages in same cycle (if no dependency hazards)
  - Superscalar: duplicate pipelines; issue multiple independent instructions per cycle
- Thread-Level: execute multiple threads on the same core concurrently
- Processor-Level: execute multiple threads/processes in parallel on multiple cores/processors

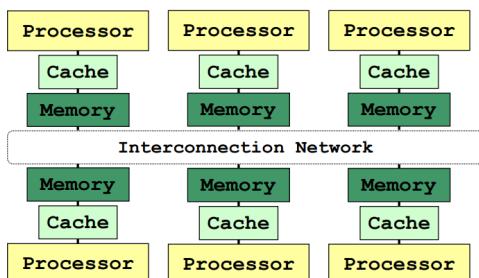
Flynn Taxonomy of Parallel Architectures:

- Single Instruction Single Data (SISD)
- Single Instruction Multiple Data (SIMD)
- Multiple Instruction Single Data (MISD)
- Multiple Instruction Multiple Data (MIMD)

## Distributed-Memory

Distributed-Memory Systems are multicenter systems:

- Node: indep. unit of processor and private memory
- Interconnection Network: supports message passing for data exchange between nodes



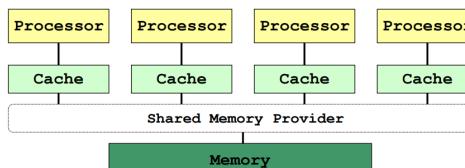
## Shared-Memory

Shared-Memory Systems are multiprocessor systems:

- Shared Memory Provider: maintains shared address space abstraction across processors
- + No need to partition code or data
- + Efficient communication: no need to physically move data between processors
- Special synchronisation constructs required
- Limited scalability due to contention

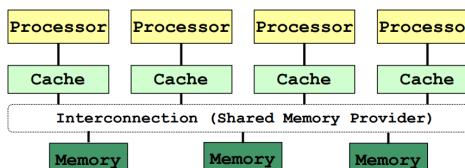
Uniform Memory Access (UMA):

- Same memory access latency
- + Suitable with few processors due to contention



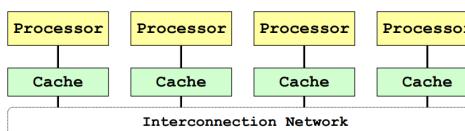
Non-Uniform Memory Access (NUMA):

- Different memory access latency for every processor (e.g. local memory faster than remote)
- Cache Coherent NUMA (ccNUMA): each node has coherent cache to reduce contention



Cache Only Memory Access (COMA):

- Memory blocks behave as caches
- Data migrates dynamically and continuously according to cache coherence scheme



## Cache Coherence

Caches reduce memory access latency and provide high bandwidth data transfer to CPU:

- Cache Size: larger cache reduces misses but increases access time (addressing complexity)
- Block Size: unit of transfer; larger blocks improve spatial locality, but longer transfer time

Write Policy:

- Write-through: write to both cache and memory
  - + Always get newest value of memory block
  - Slower due to many memory accesses; mitigate with write buffer
- Write-back: write to cache only and set dirty bit; write to memory on cache block eviction
  - + Fewer write operations to memory
  - Memory may contain invalid entries

Cache Coherence ensures all processors have consistent view of memory through their local cache. Properties:

- Program Order: if  $P$  writes to  $X$ , then n.f.w to  $X$ ,  $P$  should read same value from  $X$
- Write Propagation: if  $P_1$  writes to  $X$ , n.f.w to  $X$ ,  $P_2$  should read same value from  $X$
- Write Serialization: if  $V_1 \rightarrow X$ , then  $V_2 \rightarrow X$ , all  $P_i$  should never read  $V_2$  then  $V_1$  from  $X$

Implementations:

- Software-based: use page-fault to propagate writes
- Snooping-based: no centralized directory; caches snoop a shared bus and react to transactions
- Directory-based: centralized directory records sharing status; common in NUMA

Implications:

- "Increase" memory latency and lower cache hit rate
- Cache Ping-Pong: multiple processors repeatedly read/modifies same shared variable
- False Sharing: processors write to different addresses which map to same cache line

## Memory Consistency

Memory Consistency constrains the in which order of memory operations by a thread become visible to other threads across different memory locations.

- i. Programmer: helps reason about correctness and program behaviour
- ii. System/Compiler: decide what memory reordering is allowed (to hide write latencies)

Constraints:

- i.  $W \rightarrow R$ : write to  $X$  completes before read from  $Y$
- ii.  $R \rightarrow R$ : read from  $X$  completes before read from  $Y$
- iii.  $R \rightarrow W$ : read from  $X$  completes before write to  $Y$
- iv.  $W \rightarrow W$ : write to  $X$  completes before write to  $Y$

Sequential Consistency Model (SC):

- i. Processors issue mem. operations in program order
- ii. Result is as if all operations were interleaved in one sequential order seen consistently by all processors
- + Intuitive
- + Maintains all constraints
- Can result in loss of performance

Relaxed Consistency Model:

- i. Relax ordering if data dependencies allow (i.e. memory operations to same memory location):
  - $R \rightarrow W$ : anti-dependence (WAR)
  - $W \rightarrow W$ : output dependence (WAW)
  - $W \rightarrow R$ : flow dependence (RAW)
- ii. Allow overriding mechanism for programmers
- iii. Write-to-Read Program Order:
  - Total Store Ordering (TSO): return value written earlier without waiting for serialization
  - Processor Consistency (PC): return value of any write before propagation or serializion
- iv. Write-to-Write Program Order:
  - Partial Store Ordering (PSO): later write can propagate/serialize before earlier write
- + Hide latencies by overlapping indep. operations

## Interconnection Networks

Interconnection forms the backbone of communication between processors, memories/caches, and I/O devices.

### Topology

Topology is the geometrical shape of the connection.

Direct Interconnections (or Static, P2P) have endpoints connected directly, usually of same type (e.g. core–core). Metrics:

- i. Diameter,  $\delta(G)$ : maximum shortest-path distance between any pair of nodes
  - Small diameter  $\Rightarrow$  shorter message transmission
- ii. Degree,  $g(v)$ : no. of nodes adjacent to node  $v$ ;  $g(G)$  is the maximum node degree in network  $G$ 
  - Small degree  $\Rightarrow$  lower node hardware overhead.
- iii. Bisection Width,  $B(G)$ : minimum no. of edges removed to split the network into two equal halves
  - Bisection bandwidth  $BW(G)$ : total bandwidth between the two bisections
  - Measures capacity under simult. transmission
- iv. Node Connectivity,  $nc(G)$ : minimum no. of nodes whose failure disconnects the network
  - Determines robustness of the network
- v. Edge Connectivity,  $ec(G)$ : minimum no. of edges whose failure disconnects the network
  - Determines number of independent paths between any pair of nodes

	$n$	$g(G)$	$\delta(G)$	$ec(G)$	$B(G)$
Complete Graph	$n$	$n - 1$	1	$n - 1$	$(\frac{n}{2})^2$
Linear Array	$n$	2	$n - 1$	1	1
Ring	$n$	2	$\lfloor \frac{n}{2} \rfloor$	2	2
$d$ -dim Mesh	$r^d$	$2d$	$d(\sqrt[d]{n} - 1)$	$d$	$n^{\frac{d-1}{d}}$
$d$ -dim Torus	$r^d$	$2d$	$d\lfloor \frac{\sqrt[d]{n}}{2} \rfloor$	$2d$	$2n^{\frac{d-1}{d}}$
$k$ -dim Hypercube	$2^k$	$\log n$	$\log n$	$\log n$	$\frac{n}{2}$
$k$ -dim CCC	$k2^k$	3	$2k - 1 + \lfloor \frac{k}{2} \rfloor$	3	$\frac{n}{2k}$
Complete BT	$2^k - 1$	3	$2 \log \frac{n+1}{2}$	1	1
$k$ -ary $d$ -cube	$k^d$	$2d$	$d\lfloor \frac{k}{2} \rfloor$	$2d$	$2k^{d-1}$

Indirect Interconnection (or Dynamic) have endpoints interconnect through switches configured dynamically:

- + Reduce hardware costs by sharing switches & links

Metrics:

- i. Cost: number of switches and links
- ii. Concurrent connections

Types:

- i. Bus: set of wires carry data from sender to receiver
  - Only a pair of devices can communicate at a time
  - Bus arbiter used for coordination
  - Typically used for small number of processors
- ii. Crossbar: switch matrix (of  $n \times m$  switches) from  $n$  inputs to  $m$  outputs
  - Switch states: straight or direction change
  - Hardware costly  $\rightarrow$  small number of processors
- iii. Multistage Switching: intermediate switches with connecting wires between neighbouring stages
  - Goal: obtain small distance for any pair of input and output devices
- iv. Omega: unique path from each input to each output ( $n \times n$ )
  - Construction: switch  $(\alpha, i) \rightarrow (\beta, i+1)$  where  $\beta$  is cyclic-left-shift of  $\alpha$ , and cyclic-left-shift + inversion of LSB
  - $\log n$  stages of  $\frac{n}{2}$  switches per stage
- v. Butterfly:
  - Construction: switch  $(\alpha, i) \rightarrow (\alpha, i+1), (\alpha', i+1)$  where  $\alpha, \alpha'$  differ in the  $(i+1)^{th}$  bit from the left
- vi. Baseline:
  - Construction: switch  $(\alpha, i) \rightarrow (\beta, i+1)$  where  $\beta$  is cyclic-right-shift of last  $(k-i)$  bits of  $\alpha$ , and inversion of LSB + same cyclic-right-shift

## Routing

Routing determines path(s) from source to destination within a given topology. Algorithm Classification:

- i. Path Length:
  - Minimal: shortest-path always chosen
  - Non-minimal: shortest-path not necessary
- ii. Adaptivity:
  - Deterministic: always use same path for same pair of source and destination nodes
  - Adaptive: take into account network status (e.g. avoid congested path, avoid dead nodes)

Deterministic Algorithms:

- i. XY Routing for 2D Mesh:
  - Move in  $X$  direction until  $X_{src} == X_{dst}$
  - Move in  $Y$  direction until  $Y_{src} == Y_{dst}$
- ii. E-Cube Routing for Hypercube:
  - Bit difference in source and target node address (hamming distance) = number of hops
  - From MSB→LSB (or vice versa), find first different bit and go to neighbour node with bit corrected; at most  $n$  hops
- iii. XOR-Tag Routing for Omega Network:
  - Let  $T = \text{Source ID} \oplus \text{Destination ID}$
  - At stage  $k$ , go straight if bit  $k = 0$ , else crossover

## Additional

More Questions:

- i. Switching: how to transfer messages along a path
- ii. Flow Control: how to handle concurrent messaging

## Current Trends

Ethernet: supports P2P and broadcast; common topologies include bus, and star with a logical bus.

InfiniBand: supports P2P and multicast; common topologies include fat-tree and torus.

## 3. Programming Models

Types of Parallelism:

- i. Data Parallelism: same op. on different data elements (loop parallelism; SIMD/SPMD)
- ii. Task Parallelism: independent tasks executed concurrently (statements, loops, function calls)
  - Task Dependency Graph: DAG representing control dependency between tasks
    - Critical Path Length: slowest completion
    - Deg. Concurrency: 
$$\frac{\text{Total Work}}{\text{Critical Path Length}}$$

Representation of Parallelism:

- i. Implicit Parallelism: Haskell
- ii. Implicit Scheduling: OpenMP
- iii. Explicit Communication & Sync: MPI, Pthreads

Models of Coordination:

- i. Shared Address Space:
  - Communicate via shared variables; mutual exclusion via locks
  - Matches shared-memory (UMA/NUMA); requires HW support; costly to scale
- ii. Data Parallel:
  - Map a function onto a large collection of data (historic: SIMD/vector; modern: CUDA)
  - Rigid structure enables parallel scheduling; ideally no communication among invocations
- iii. Message Passing:
  - Private address spaces; explicit send/receive
  - Matches distributed memory
  - Point-to-point vs Global communication
  - Blocking vs Non-Blocking
  - Buffered vs Non-Buffered
  - Synchronous vs Asynchronous

## Foster's Design Methodology

- i. Partitioning: problem into pieces/data
  - $\geq 10\times$  primitive tasks than cores, of similar size
  - Minimize redundant computations and data
  - Number of tasks increases with problem size
- ii. Communication: pass needed data among tasks
  - Balance communication among tasks
  - Perform communication in parallel
  - Overlap computation with communication
  - Communicate with small number of neighbours
- iii. Agglomeration: combine tasks into larger tasks
  - Number of tasks  $\geq$  number of cores
  - Reduce cost of task creation + communication
  - Maintain scalability, simplify code
- iv. Mapping: tasks to execution units
  - Maximize processor utilization; minimize IPC
  - Optimal mapping is NP-hard  $\Rightarrow$  heuristics
  - Consider static vs dynamic allocation;  
if static, tasks:cores  $\geq 10:1$ ;  
if dynamic, allocator must not bottleneck

## Parallel Programming Patterns

Patterns:

- i. Fork-Join: `fork` child tasks, then `join`
- ii. Parbegin-Parend: mark statements to execute in parallel, then sync at end
- iii. SIMD: single instruction stream over multiple data elements (synchronous)
- iv. SPMD: same program on multiple threads; different data via IDs (asynchronous)
- v. Master-Worker: master distributes tasks to workers
- vi. Client-Server (MPMD): request-response model
- vii. Task Pools: shared pool of tasks; workers dynamically dequeue tasks
- viii. Producer-Consumer: produce tasks into buffer for consumers to retrieve; needs sync
- ix. Pipelining: split compute into stages for data stream; throughput improves via stage-overlap

## OpenMP

Shared-Memory Programming Model:

- i. Threads communicate via shared variables; may also have private variables
- ii. Avoid race conditions using mutual exclusion

### Parallel For (Data Parallel)

```
// Parallel outer loop; shared data +
private indices
#pragma omp parallel for num_threads(8)
shared(a,b,result) private(i,j,k)
for (i = 0; i < size; i++)
    for (j = 0; j < size; j++)
        for (k = 0; k < size; k++)
            result[i][j] += a[i][k] * b[k][j];
```

### Mutex with OpenMP Lock

```
int count = 0;
omp_lock_t lock;
omp_init_lock(&lock);

#pragma omp parallel
{
    omp_set_lock(&lock);
    count = count + 1;
    omp_unset_lock(&lock);
}
```

## NVIDIA CUDA

CUDA Programming Model:

- i. Heterogeneous: Host (CPU) + Device (GPU)
- ii. Kernel: function that runs on the device
- iii. Thread: local registers and memory; `threadIdx`
- iv. Warp: execution unit of 32 threads scheduled together (SIMT); one instruction at a time
- v. SM (Streaming Multiprocessor): schedules warps with resources (register file, shared memory); thread blocks execute concurrently on an SM

Function Qualifiers:

Qualifier	Scope	Usage
<code>__host__</code>	CPU	local to host
<code>__device__</code>	GPU	local to device
<code>__global__</code>	GPU	launch from host

Variable Qualifiers:

Qualifier	R/W	Cache	Scope	Lifetime
<code>__device__</code>	R/W	No	grid	static
<code>__constant__</code>	R	Yes	grid	static
<code>__shared__</code>	R/W	No	block	block
unqualified	R/W	—	thread	thread
<code>__managed__</code>	R/W	—	global	static

Optimizations:

- i. Reduce transfers: minimize host↔device copies; batch small copies; use pinned memory; overlap copy/compute with `cudaMemcpyAsync` + streams
- ii. Global memory: coalesce accesses; minimize global loads/stores; use `__shared__`; avoid bank conflicts
- iii. Occupancy: many warps to hide latency; threads/block multiple of 32; balance registers/shared-mem so  $\geq 1$  block/SM
- iv. Control flow: minimize warp divergence (avoid thread-ID dependent branches)
- v. Compute: prefer high-throughput ops; single precision faster; avoid int div/mod (use shifts/bitwise)
- vi. Avoid multi-context: do not create multiple CUDA contexts per GPU within one application

## Matrix Addition (Device)

```
--global--
void addMatrixG(float *a, float *b, float *c, int N)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;
    int index = i + j * N;
    if (i < N && j < N)
        c[index] = a[index] + b[index];
}
```

## Matrix Multiplication (Device)

```
--global--
void addMatrixG(float *a, float *b, float *c, int N)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;
    int index = i + j * N;
    if (i < N && j < N)
        c[index] = a[index] + b[index];
}
```

## Host Skeleton

```
void main(){
    // Assume host arrays: h_A, h_B, h_C; device
    // pointers: d_A, d_B, d_C
    size_t bytes = (size_t)N * sizeof(float);
    cudaMalloc((void**)&d_A, bytes);
    cudaMalloc((void**)&d_B, bytes);
    cudaMalloc((void**)&d_C, bytes);

    // Copy inputs Host -> Device
    cudaMemcpy(d_A, h_A, bytes,
               cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, bytes,
               cudaMemcpyHostToDevice);

    // Configure grid/block and launch kernel
    dim3 dimBlk(BLOCK_X, BLOCK_Y);      // e.g.
                                         (16,16)
    dim3 dimGrd((NX + BLOCK_X - 1)/BLOCK_X,
                (NY + BLOCK_Y - 1)/BLOCK_Y);
    kernel<<<dimGrd, dimBlk>>>(d_A, d_B, d_C, ...);

    // Copy result Device -> Host
    cudaMemcpy(h_C, d_C, bytes,
               cudaMemcpyDeviceToHost);

    // Free device memory
    cudaFree(d_A); cudaFree(d_B); cudaFree(d_C);
}
```

# MPI

Message Passing Model:

- i. Distributed memory
- ii. Explicitly represent parallelism with send/receive

MPI Program Overview:

- i. Initialize communications
- ii. Communicate to coordinate computation
- iii. Exit from message-passing system

## Communication Semantics

Protocol Choice:

- i. Buffered: sender returns after copying data into a communication buffer (less idling; buffer overhead)
- ii. Non-buffered: sender blocks until matching receive is encountered (idling + deadlock risks if mismatch)

Local View:

- i. Blocking: return  $\Rightarrow$  safe to reuse resources/buffers
- ii. Non-blocking: may return before completion; programmer must ensure completion (poll/check)

Global View:

- i. Synchronous: operation does not complete before both sides have started
- ii. Asynchronous: sender can proceed without coordination with receiver

Operations:

- i. Sync + Blocking: MPI\_Ssend
- ii. Async + Blocking: MPI\_Send / MPI\_Recv
- iii. Sync + Non-blocking: MPI\_Issend
- iv. Async + Non-blocking: MPI\_Isend / MPI\_Irecv

\* Note: blocking and non-blocking ops can be mixed

## Library Functions

Initi, Finalize, Abort Calls:

- i. `int MPI_Init(int* argc, char** argv[])`
- ii. `int MPI_Finalize(void)`
- iii. `int MPI_Abort(MPI_Comm comm, int errorCode)`

Point-to-Point Messaging Calls:

- i. `int MPI_Send(void* buf, int count, MPI_Datatype dt, int dst, int tag, MPI_Comm c)`
- ii. `int MPI_Recv(void* buf, int count, MPI_Datatype dt, int src, int tag, MPI_Comm c, MPI_Status *status)`

Message Format:

- i. Data: start-buffer + count + datatype
- ii. Envelope: destination/source (rank) + tag + communicator
- iii. Receive buffer must be  $\geq$  message length
- iv. Wildcards: MPI\_ANY\_SOURCE, MPI\_ANY\_TAG
- v. `MPI_Status`: MPI\_SOURCE, MPI\_TAG, MPI\_ERROR

Ordering:

- i. 1 sender, 1 receiver: messages delivered in order
- ii.  $> 2$  processes: message delivery order undefined

## Deadlocks

Deadlock Types:

- i. Message Order: both sides wait on receives
- ii. Buffer: no/small buffers  $\Rightarrow$  sends cannot complete

Secure MPI program: correctness does not depend on assumptions about MPI runtime properties.

- Specify send/receive order  
(e.g. even: send $\rightarrow$ recv, odd: recv $\rightarrow$ send)

## Collective Communication

Process Groups are ordered sets of processes:

- i. Unique rank identifier
- ii. Processes may be in multiple groups
- iii. Communicator: communication domain for a group
  - Intra-communicator: collectives in a group; default MPI\_COMM\_WORLD
  - Inter-communicator: P2P between two groups

Virtual Topology is a Cartesian-style communicator:

- i. `MPI_Cart_create, MPI_Cart_get, MPI_Cartdim_get, MPI_Cart_coords, MPI_Cart_rank, MPI_Cart_shift`

Collectives involve all processes in a communicator:

- i. Total Exchange: `MPI_Alltoall`
- ii. Multi-broadcast: `MPI_Allgather`
- iii. Multi-accumulation: `MPI_Reduce_scatter`
- iv. Scatter/Gather: `MPI_Scatter / MPI_Gather`
- v. Single Broadcast: `MPI_Broadcast`
- vi. Single Accumulation: `MPI_Reduce`
- vii. Barrier: `MPI_Barrier`

## Hello World (Master Sends)

```
MPI_Init(&argc, &argv);
MPI_Comm_size(MPI_COMM_WORLD, &size);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
tag = 100;
if (rank == 0) {
    strcpy(message, "Hello World 2");
    for (i=1; i<size; i++)
        MPI_Send(message, 14, MPI_CHAR, i, tag,
                 MPI_COMM_WORLD);
} else {
    MPI_Recv(message, 14, MPI_CHAR, 0, tag,
             MPI_COMM_WORLD, &status);
}
printf("node %d : %.13s\n", rank, message);
MPI_Finalize();
```

## 4. Performance

Performance is measured by primary goals:

- i. Response Time: wall-clock duration of an execution
  - User CPU Time: time spent on program
  - System CPU Time: time spent on OS routines
  - Waiting Time: time spent waiting for I/O and execution of other programs (timesharing)
- ii. Throughput: average work completed per unit time

### Sequential Programs

User CPU Time:

$$T_{user} = N_{cycle} \times T_{cycle}$$

- i.  $T_{cycle} = \frac{1}{\text{clock rate}}$
- ii.  $N_{cycle} = \sum n_i \times CPI_i \approx N_{instr} \times CPI$

User CPU Time (Memory Access; One-level Cache):

$$T_{user} = (N_{cycle} + N_{mm\_cycle}) \times T_{cycle}$$

- i.  $N_{mm\_cycle} = N_{read\_cycle} + N_{write\_cycle}$
- ii.  $N_{read\_cycle} = N_{reads} \times R_{read\_miss} \times N_{miss\_cycles}$

Average Memory Access Time (Two-level Cache):

$$T_{read} = T_{hit} + R_{miss} \times T_{miss}$$

- i.  $T_{read} = T_{hit}^{L1} + R_{miss}^{L1} \times T_{miss}^{L1}$
- ii.  $T_{miss}^{L1} = T_{hit}^{L2} + R_{miss}^{L2} \times T_{miss}^{L2}$
- iii. Global miss rate:  $R_{miss}^{L1} \times R_{miss}^{L2}$

Throughput:

$$\text{i. } MIPS = \frac{N_{instr}}{T_{user} \times 10^6} = \frac{\text{clock\_freq}}{CPI \times 10^6}$$

$$\text{ii. } MFLOPS = \frac{N_{fl\_ops}}{T_{user} \times 10^6}$$

### Parallel Programs

Parallel Execution Time on  $p$  processors,  $T_p$  consists:

- i. Time for executing local computations
- ii. Time for exchange of data between processors
- iii. Time for synchronization between processors
- iv. Waiting time

Cost:

$$C_p = p \times T_p$$

- i. Measures total work performed by all processors
- ii. Cost-optimal: parallel program executes same total number of operations as fastest sequential program

Speedup:

$$S_p = \frac{T_*}{T_p}$$

- i. Measures the benefit of parallelism compared to best sequential algorithm time  $T_*$
- ii. Theoretically,  $S_p \leq p$  always holds
- iii. In practice,  $S_p > p$  (superlinear speedup) can occur (e.g. problem working set “fits” in the cache)

Efficiency:

$$E_p = \frac{T_*}{C_p} = \frac{S_p}{p} = \frac{T_*}{p \times T_p}$$

- i. Measures actual degree of speedup performance achieved compared to maximum
- ii. Ideal speedup  $S_p = p \Rightarrow E_p = 1$

### Scalability

Scalability is the interaction between problem size and parallel machine size:

- i. Impacts: load balancing, overhead, arithmetic intensity, locality (application dependent)
- ii. Fixed problem size:
  - Small  $n$ : overheads dominate  $\Rightarrow$  suitable for small machine, poor on large machine
  - Large  $n$ : working set may not fit on small machine (thrashing / exceeds cache / can't run)
- iii. Constraints:
  - Application: particles/processor (N-body), transactions/processor (distributed DB)
  - Resources: time-limited, memory-limited, problem-limited (same problem faster)

Amdahl's Law:

$$S_p(n) = \frac{T_*(n)}{f \cdot T_*(n) + \frac{1-f}{p} T_*(n)} = \frac{1}{f + \frac{1-f}{p}} \leq \frac{1}{f}$$

- i. Sequential fraction:  $0 \leq f \leq 1$  limits speedup
- ii. Discourage from making large parallel computers, focus on reducing sequential fraction  $f$
- iii. Rebuttal: circumvented for large problem size
  - Commonly  $f$  is a function of  $n$ ,  $f(n)$
  - Efficient parallel algorithm:  $\lim_{n \rightarrow \infty} f(n) = 0$
  - Thus,  $\lim_{n \rightarrow \infty} S_p = \frac{p}{1+(p-1)f(n)} = p$

Gustafson's Law:

$$S_p(n) = \frac{\tau_f + \tau_v(n, 1)}{\tau_f + \tau_v(n, p)}$$

- i. Encouraged to use more processors to solve larger problem in the same time (in CPU-bound apps)
- ii.  $\tau_f$ : constant sequential execution time
- iii.  $\tau_v(n, p)$ : parallel execution time on  $p$  processors; if perfectly parallelizable (no overheads):
  - $\tau_v(n, 1) = T^*(n) - \tau_f$
  - $\tau_v(n, p) = \frac{T^*(n) - \tau_f}{p}$
- iv. If  $T^*(n)$  increases strongly monotonically with  $n$ , then  $\lim_{n \rightarrow \infty} S_p(n) = p$

## Communication

Communication Performance Measures:

- i. Bandwidth,  $B$ : maximum data rate (bytes/s)
- ii. Throughput: effective bandwidth (bytes/s)
- iii. Time of Flight,  $T_{delay}$ : time for first bit to arrive
- iv. Byte Transfer Time,  $t_B$ : time to transmit one byte
- v. Transmission Time: time to transmit message
- vi. Transport Latency: time to transfer a message (time of flight + transmission time)
- vii. Sender Overhead,  $O_{send}$ : time to compute checksum, append header, and execute routing
- viii. Receiver Overhead,  $O_{recv}$ : time to compare checksum and generate acknowledgement

Communication Time (message size  $m$ ):

$$\begin{aligned} T(m) &= O_{send} + T_{delay} + \frac{m}{B} + O_{recv} \\ &= T_{overhead} + \frac{m}{B} \\ &= T_{overhead} + t_B \cdot m \end{aligned}$$

where:

- i.  $T_{overhead} = O_{send} + T_{delay} + O_{recv}$  is independent of message size  $m$
- ii. Byte Transfer Time,  $t_B = \frac{1}{B}$
- iii. Assume no checksum error, and no network contention / congestion

## Performance Analysis

Practical Workflow:

- i. Measure first; define target metric (latency vs throughput vs speedup)
- ii. Localise the bottleneck (CPU vs memory vs sync vs I/O vs network)
- iii. Change *one* factor at a time; keep workload and environment fixed
- iv. Validate: speedup should be explained by reduced dominant component

## Instrumentation

Performance Instrumentation finds performance metrics:

- i. Latency: time to service a request
- ii. Response Time: wall-clock duration of an execution
- iii. Throughput: average work completed per unit time
- iv. IOPS: I/O operations per second
- v. Utilization: fraction of time a resource is busy
- vi. Saturation: degree to which a resource has queued work it cannot service

Perspectives:

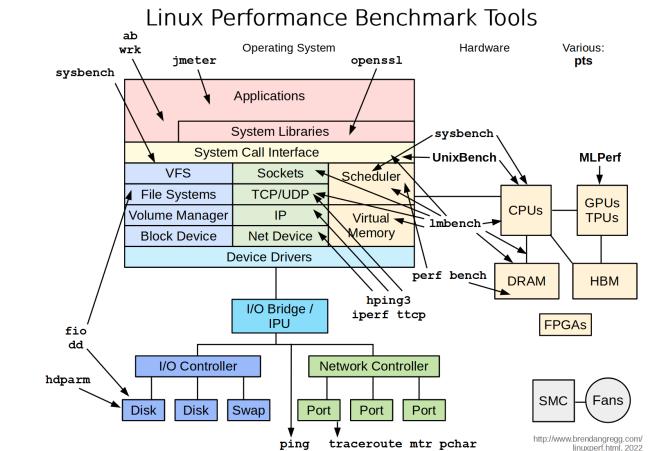
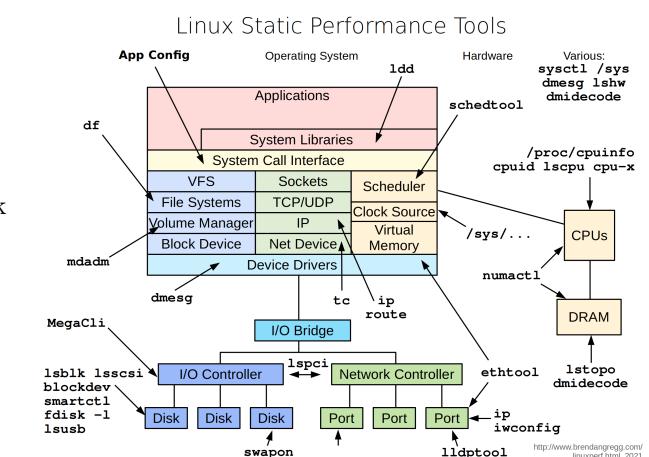
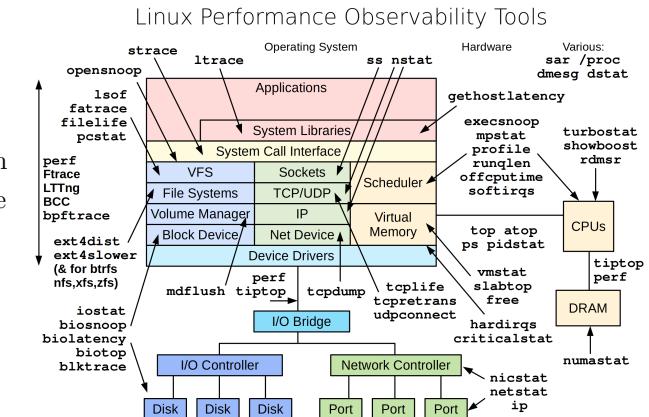
- i. Resource Analysis: focus on utilization
  - Time-based: average time the resource was busy
  - Capacity-based: resource throughput
- ii. Workload Analysis: focus on throughput, latency

Methodologies:

- i. Anti-Methodologies: undeliberate, street light (look for obvious issues), drunk man (tune at random)
- ii. Problem Statement
- iii. USE (per resource): Utilization, Saturation, Errors
- iv. Monitoring: record metrics over time

Tool Types & Counters:

- i. Observability: watch activity under workload (timing statements, performance counters)
- ii. Static: examine system at rest
- iii. Benchmarking: load test (caution: contention/production risk)
- iv. Tuning: change defaults (danger: regressions now/later under load)
- v. Categorization:
  - System-wide vs per-process
  - Fixed counters (kernel-maintained metrics) vs event-based counters (enabled as needed)
  - Profiling: samples/snapshots vs Tracing: record every event occurrence



## 5. Energy Efficiency

Energy-Efficiency Motivation:

- i. Power dissipation has become a limiting factor (“power wall”)
- ii. Higher performance  $\Rightarrow$  more/faster computers  $\Rightarrow$  more power/heat/cooling/space/cost

## Mobile Computing

Trends:

- i. Energy-efficiency: decrease power consumption in hardware and increase performance
- ii. Increase battery capacity
- iii. ARM processing systems:
  - Designed for power-efficiency
  - Closing in on x64 performance
  - Highly customizable architecture

ARM big.LITTLE (low-power approach):

- i. Big CPU: high performance for intensive workloads
- ii. Little CPU: low power for majority workloads
- iii. Switch between cores depending on demand

## Heterogeneous Computing

Heterogeneous Platforms:

- i. Processors: brawny vs wimpy, big-little
- ii. Supercomputers: accelerators
- iii. Data Centers: mixed server generations
- iv. Cloud: heterogeneous resources with different price-performance

Challenges:

- i. Use heterogeneity to reduce power while maintaining performance
- ii. Energy-efficient configuration for parallel applications is complex (scheduling)
- iii. Burden shifts to programmers for portable code

Data Centers:

- i. Large-scale data centers consume significant electricity and emit substantial CO<sub>2</sub>
- ii. Cooling overhead can be comparable to computation energy
- iii. Power Use Efficiency (PUE) =  $\frac{\text{Total Energy}}{\text{IT Energy}}$ : measures energy efficiency in data centers

Google (energy-efficiency practices):

- i. Continuously measure efficiency:
  - Computation (IT) energy
  - Overhead: cooling and conversions
- ii. Build custom, highly-efficient servers:
  - Minimize AC/DC conversion losses
  - Remove unnecessary parts
  - Strategic rack placement
  - Tune fan speeds for cooling
- iii. Extend equipment lifecycle
  - Reuse, resell components
- iv. Control equipment temperature:
  - Raise operating temperature (e.g. 26°C)
  - Thermal modeling
- v. Cooling with water instead of chillers:
  - Evaporative cooling, sea water, recycled water

Reducing Energy Consumption:

- i. Move less data:
  - Reduce transfers to/from memory
  - Exploit locality
  - Use compression
- ii. Use specialized processing:
  - Compute less (no parallel if it increases work)
  - CPU-like cores + GPU-like throughput-optimized cores
  - FPGAs (programmable hardware)

## 6. Appendix

name	prefix	multiplier
exa	E	$10^{18}$
peta	P	$10^{15}$
tera	T	$10^{12}$
giga	G	$10^9$
mega	M	$10^6$
kilo	K	$10^3$
milli	m	$10^{-3}$
micro	$\mu$	$10^{-6}$
nano	n	$10^{-9}$
pico	p	$10^{-12}$

name	prefix	multiplier
exbi	Ei	$2^{60}$
pebi	Pi	$2^{50}$
tebi	Ti	$2^{40}$
gibi	Gi	$2^{30}$
mebi	Mi	$2^{20}$
kibi	Ki	$2^{10}$