


CS3210 Tutorial 3


AY 25/26 Sem 2 — [github/omgeta](https://github.com/omgeta)

- Q1. Instruction-Level Parallelism: pipeline, branch prediction, superscalar
Thread-Level Parallelism: Simultaneous Multi-Threading (SMT)
Processor-Level Parallelism: multi-core running in parallel, UPI shared memory, Omni-Path distributed memory
IO Parallelism: Multiple DDR4 channels and PCIe lanes
- Q2. (a) SISD: single core with single instruction stream on each memory address
(b) MIMD: parallel instruction streams on different cores can act on different data
(c) SIMD: vector instructions act over multiple memory locations
(d) SISD: single core with single instruction stream, pipelined instructions still act on a single memory address at a specific time step
(e) MISD: same paper (data) operated on by different students (instructions)
- Q3. (a) False; Shared-Memory can be UMA or NUMA
(b) False; data locality affects latency
(c) True
- Q4. (a) False/True; binary semaphores emulate mutual exclusion but code depending on mutex ownership and semantics may be affected
(b) False; it depends on overheads, synchronisation and workload


Q5. (a) Fragment 1:




Fragment 2:



(b) Fragment 1:




Max Concurrency = 4, Average Concurrency = $\frac{59}{28} = 2.1$
Fragment 2:




Max Concurrency = 4, Average Concurrency = $\frac{59}{33} = 1.78$

(c) Fragment 1 Speedup = $\frac{59}{28} = 2.1$
Fragment 2 Speedup = $\frac{59}{33} = 1.78$

(d) Fragment 1:



Speedup = $\frac{59}{32} = 1.84 < 2.1$
Fragment 2:



Speedup = $\frac{59}{35} = 1.69 < 1.78$

- Q6. (a) Data Parallelism; SIMD; each independent output element c_i is an independent dot product of the vector b with row a_i
- (b) Parbegin-Parend
- Q7. (a) Task Parallelism: each stage is a task
 Pipelining: best for cases where each stage takes a similar time
- (b) Task Parallelism: each stage is a task
 Producer–Consumer: producer reads from socket and processes, consumer writes back to socket and there are equal number of producers as consumers
- (c) Task Parallelism: each stage is a task
 Producer–Consumer: same as (b) but more producers than consumers
- (d) Task Parallelism: each stage is a task
 Task Pool: pool of tasks are ready for jobs and assigned as necessary
- (e) Same; but need explicit communication

Q8. Average CPI (Translation 1) = $\frac{10}{5} = 2$
 Average CPI (Translation 1) = $\frac{9}{6} = 1.5$

Q9. Execution Time (A_1) = $\frac{(5+1\cdot 2+1\cdot 3)\times 10^9}{2\times 10^9} = 5s$
 Execution Time (A_1) = $\frac{(10+1\cdot 2+1\cdot 3)\times 10^9}{2\times 10^9} = 7.5s$

Conclusion: MIPS is not an accurate measurement of performance

- Q10. (a) Sequential Time: $\frac{(100+100^2)\times 2}{10^9} = 20200ns$
 Parallel Time ($p = 10$): $\frac{(100+\frac{100}{10})\times 2}{10^9} = 2200ns$ ($9.18\times$ speedup)
 Parallel Time ($p = 100$): $\frac{(100+\frac{100}{100})\times 2}{10^9} = 400ns$ ($50.6\times$ speedup)
 Parallel Time ($p = \infty$): $\frac{(100+0)\times 2}{10^9} = 200ns$ ($101\times$ speedup)
- (b) From (a), $T_1 = 20200ns$
 For $p = 10$,

$$\frac{(N + \frac{N_2}{10}) \times 2}{10^9} = 20200ns \implies N_{p=10} = 312 \implies S_{10}(312) = 9.669\times$$

 For $p = 100$,

$$\frac{(N + \frac{N_2}{100}) \times 2}{10^9} = 20200ns \implies N_{p=100} = 956 \implies S_{100}(312) = 90.58\times$$